

The CLR and IL in a Nutshell

Bart J.F. De Smet
bartde@outlook.com



pluralsight 
hardcore dev and IT training

A Whirlwind Tour...

- **Quick refresher of:**
 - CLR internals
 - IL instructions

- **Check out our course catalog!**
 - CLR Fundamentals, by Mike Woodring
 - MSIL for the C# Developer, by Filip Ekberg

- **Books**
 - CLR via C#, Jeffrey Richter
 - Expert .NET 2.0 IL Assembler, Serge Lidin
 - Shared Source CLI 2.0 Internals, Joel Poebar & Ted Neward


Quick Intro to IL Code

- **Intermediate Language (IL)**
 - Virtual Machine language of the CLR
 - Emitted by managed language compilers (C#, VB, F#, etc.)
 - NGEN or JIT compiled to native code
 - Stack-based evaluation
 - No registers
 - Locals, arguments, fields, etc.
 - Verification
 - Type safe
 - Memory safe
 - Leverages metadata

Quick Intro to IL Code

■ Stack-based evaluation

- “Scratch-pad” for computation
- Pop *operands*, execute *operator*, push result
- Each instruction has a net effect on the stack
 - E.g. `add` = 2 x pop + 1 x push
- Stack transitions for “1 + 2”



Quick Intro to IL Code

■ Essential instructions

- pop
 - Pops the object on top of the stack
 - Often used to discard stuff to rebalance the stack

- dup
 - Duplicates the object on top of the stack
 - Often used to eliminate loads and stores to locals

- nop
 - No-operation, doesn't do anything "useful"
 - Often used in non-optimized builds (csc /o-)
 - Breakpoints can only be set on instructions
 - E.g. emit nop for lines with curly braces

Quick Intro to IL Code


■ Loading constants

- Numerical:
 - ldc.i4 for “load constant integer 4 bytes” (int)
 - ldc.r8 for “load constant floating point 8 bytes” (double)
 - Value as an operand in the IL instruction stream
 - Shorthand instructions, e.g. ldc.i4.1 for Int32 1, ldc.i4.m1 for Int32 -1
- Boolean:
 - Represented as 0 (false) or 1 (true)
- String:
 - ldstr for “load string”
 - Value as an operand that points to a string table entry
- Null:
 - ldnull
 - Null reference, useful for initialization or cleanup of locations

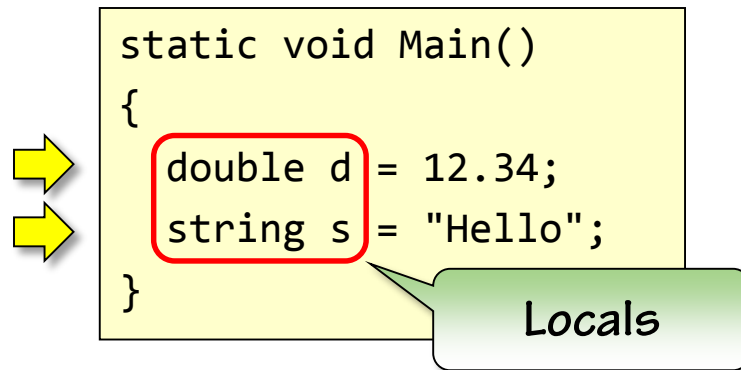
Quick Intro to IL Code

Locals

- Typed slots to hold objects
- JIT may put those in machine registers or on the stack
- Instructions to load and store
 - E.g. `ldloc.0` pushes the 0th local onto the evaluation stack
 - E.g. `stloc.1` pops from the evaluation stack and stores into the 1st local



Quick Intro to IL Code



```
.method private hidebysig static void Main() cil managed  
{  
    .entrypoint  
    .maxstack 1  
    .locals /*11000001*/ init (float64 V_0, string V_1)  
    IL_0000: /* 00 | */ nop  
    IL_0001: /* 23 | AE47E17A14AE2840 */ ldc.r8 12.34  
    IL_000a: /* 0A | */ stloc.0  
    IL_000b: /* 72 | (70)000001 */ ldstr "Hello" /* 70000001 */  
    IL_0010: /* 0B | */ stloc.1  
    IL_0011: /* 2A | */ ret  
}
```

Op codes

Quick Intro to IL Code

- **Arithmetic, relational, logical, conversions**
 - Arithmetic instructions
 - neg, add, sub, mul, div, rem, shl, shr
 - Overflow variants with .ovf suffix (cf. checked in C#)
 - Relational instructions
 - clt, cle, cgt, cge, ceq, cne (also branch counterparts)
 - Push zero or one based on outcome of comparison
 - Bitwise logical instructions
 - and, or, xor, not
 - Short-circuiting behavior requires control flow
 - Conversions
 - conv.<type>, e.g. conv.i4 to convert to 4-byte integer
 - isinst for type checks
 - castclass for casts to a specified type

Quick Intro to IL Code

■ Branch instructions

- Unconditional branch instructions
 - br – method-local branch using the specified offset
- Conditional branch instructions
 - brtrue, brfalse – check top of stack for true or false
 - beq, bne, blt, ble, bgt, bge – relational operators with branching
- Variants
 - .s suffix – short branch if offset fits in 1 byte
 - .un – unsigned variants
- Switch tables
 - switch – jumps based on integer operand

Quick Intro to IL Code

Branch

```
static void Gt(int a, int b) {  
    if (a > b)  
        Console.WriteLine("a > b");  
    else  
        Console.WriteLine("a <= b");  
}
```

Fall-through

```
.method private hidebysig static void Gt(int32 a, int32 b) cil managed  
{  
    IL_0000: ldarg.0  
    IL_0001: ldarg.1  
    IL_0002: ble.s      IL_000f      /* 31 (ble.s) | 0B (offset) */  
    IL_0004: ldstr      "a > b"  
    IL_0009: call       void [mscorlib]System.Console::WriteLine(string)  
    IL_000e: ret  
    IL_000f: ldstr      "a <= b"  
    IL_0014: call       void [mscorlib]System.Console::WriteLine(string)  
    IL_0019: ret  
}
```

Short branch

Quick Intro to IL Code

■ Call stacks

- Different call instructions
 - call – regular “direct” call
 - callvirt – call with virtual dispatch
 - calli – indirect call through a pointer (interop)
- Arguments
 - Push arguments to make a method call
 - Arguments passed left-to-right on the stack
 - Instance methods hold “this” in 0th argument
 - Call stack frames hold locals and arguments
 - Instruction to load an argument: ldarg
- Return to caller using ret instruction
 - Stack should only contain one object (or none if void)
- Exceptions unwind the call stack

Quick Intro to IL Code

```
static void Main()
{
    double d = Math.Pow(3.14, 2.81);
    Console.WriteLine(d);
}
```

```
.method private hidebysig static void  Main() cil managed
{
    .entrypoint
    .maxstack 2
    .locals init (float64 V_0)
    IL_0000:  ldc.r8      3.140000000000000001
    IL_0009:  ldc.r8      2.810000000000000001
    IL_0012:  call       float64 [mscorlib]System.Math::Pow(float64, float64)
    IL_0017:  stloc.0
    IL_0018:  ldloc.0
    IL_0019:  call       void [mscorlib]System.Console::WriteLine(float64)
    IL_001e:  ret
}
```

Quick Intro to IL Code

■ Calls in C#

- Static methods use call
- Instance methods use callvirt
 - Even for non-virtuals
 - Performs null check for v-table

```
.class Foo {  
    .method void Bar(int32 a) {  
        ldarg.1  
        call void Console::WriteLine(int32)  
        ret  
    }  
}
```

Does not use
ldarg.0

■ Special calls

- Use native implementation provided by the CLR
- [MethodImpl(MethodImplOptions.InternalCall)]
 - Often used for reflection-related stuff
 - E.g. System.Object::GetType
- [DllImport("QCall")]
 - Quick calls
 - from the mscorlib.dll assembly
 - to native helper methods in mscorwks.dll or clr.dll
 - E.g. System.GC::_Collect

Quick Intro to IL Code

■ Throwing exceptions

- throw
 - Used for “throw ex;” in C#
 - Resets the stack trace
 - ExceptionDispatchInfo in .NET 4.5
- rethrow
 - Used for “throw;” in C#

```
try {  
    // Do stuff  
}  
catch (SomeException ex) {  
    throw; // rethrow  
}
```

■ Handling exceptions

- Metadata describes protected regions, i.e. “try”...
 - “catch” – type-based exception handling
 - “finally” – cleanup upon successful or exceptional exit
 - “fault” – not available in C#
 - “filter” – conditional exception handling available in Visual Basic
- Control flow in and out of a protected region
 - leave instruction to exit a protected region
 - causes handlers to run
 - endfinally and endfault to exit handlers

Quick Intro to IL Code

■ Working with objects

- newobj – creates a new object on the GC heap
 - Causes memory allocation and can throw OutOfMemoryException
 - Runs specified constructor on object with zeroed memory
 - Returns reference to newly created object
- ldfld and stfld
 - Loads and stores from/to fields
 - Parameterized by metadata token of the fields

■ Arrays (one-dimensional)

- newarr – creates a new array of the specified length
- ldlen – loads the array length
- ldelem and stelem – loads and stores array elements using an index


JIT and NGEN

- **Compilation model**

- Front-end

- Managed language compilers, such as C#, Visual Basic, F#
 - Emit IL code

```
static int Add(int a, int b)
{
    return a + b;
}
```



```
.method static int32 Add(
    int32 a, int32 b)
{
    ldarg.0
    ldarg.1
    add
    ret
}
```

JIT and NGEN

■ Compilation model

- Front-end
 - Managed language compilers, such as C#, Visual Basic, F#
 - Emit IL code
- Back-end
 - Just in Time (JIT) compilation to x86, x64, or ARM at runtime, or
 - Native Image Generation (NGEN) ahead of time, e.g. during setup

```
7ff8`1c8100d0  mov  dword ptr [rsp+10h],edx
7ff8`1c8100d4  mov  dword ptr [rsp+8],ecx
7ff8`1c8100d8  mov  ecx,dword ptr [rsp+10h]
7ff8`1c8100dc  mov  eax,dword ptr [rsp+8]
7ff8`1c8100e0  add  eax,ecx
7ff8`1c8100e2  jmp  00007ff8`1c8100e4
7ff8`1c8100e4  nop
7ff8`1c8100e5  ret
```



JIT

```
.method static int32 Add(
    int32 a, int32 b)
{
    ldarg.0
    ldarg.1
    add
    ret
}
```

Intermezzo: Compiler Tradeoffs

Front-end (C#)



- ☐ Developer productivity
- ☐ Global program knowledge
- ☐ Machine agnostic
- ☐ More time to optimize
- ☐ Can defer to JIT

Back-end (JIT)



- ☐ Efficient execution
- ☐ Local program knowledge
- ☐ Machine knowledge
- ☐ Less time to optimize
- ☐ Last in line

JIT and NGEN

- **Using SOS to analyze a method**
 - Method descriptor (“md”)
 - Internal identifier of the method
 - !name2ee <module> <method>
 - Display method info
 - Class, method table, JIT status, etc.
 - !dumpmd <method desc>
 - Break on method
 - Native instructions only appear after JIT
 - !bpmd command
 - !bpmd -md <method desc>
 - !bpmd <module> <method>
 - !bpmd <source file>:<line> (if PDBs are available)
 - Show code
 - !dumpil <method desc>
 - !U <method desc>

JIT and NGEN

```
C:\Demo> cdb.exe arith.exe
```

```
0:000> !bpmd arith.exe Arith.Main
```

```
Found 1 methods in module 00007ff81c6f2fb0...
```

```
MethodDesc = 00007ff81c6f3fe8
```

```
Adding pending breakpoints...
```

```
0:000> g
```

```
(44c.1934): CLR notification exception - code e0444143 (first chance)
```

```
JITTED arith!Arith.Main()
```

```
Setting breakpoint: bp 00007FF81C810090 [Arith.Main()]
```

```
Breakpoint 0 hit
```

```
0:000> !u 00007ff81c810090
```

```
Normal JIT generated code
```

```
Arith.Main()
```

```
7ff8`1c810090  sub    rsp,28h
```

```
7ff8`1c810094  mov    edx,2
```

```
7ff8`1c810099  mov    ecx,1
```

```
7ff8`1c81009e  call   7ff8`1c6fc028 (Arith.Add(Int32, Int32), mdToken: 000006000002)
```

```
7ff8`1c8100a3  mov    ecx,eax
```

```
7ff8`1c8100a5  call   mscorlib_ni+0xd24780 (7ff8`7b5d4780) (Console.WriteLine(Int32),
```



JIT thunk

JIT and NGEN

Invoke JIT

```
0:000> u 00007ff8`1c6fc028
00007ff8`1c6fc028 e87365755f call    clr+0x25a0 (00007ff8`7be525a0)
00007ff8`1c6fc02d 5e      pop     rsi
00007ff8`1c6fc02e 0201    add     al,byte ptr [rcx]
```

```
0:000> !bpmd arith.exe Arith.Add
```

```
Found 1 methods in module 00007ff81c6f2fb0...
```

```
MethodDesc = 00007ff81c713ff8
```

```
Adding pending breakpoints...
```

```
0:000> g
```

```
(44c.1934): CLR notification exception - code e0444143 (first chance)
```

```
JITTED arith!Arith.Add(Int32, Int32)
```

```
Setting breakpoint: bp 00007FF81C8100E3 [Arith.Add(Int32, Int32)]
```

```
Breakpoint 1 hit
```

JITted code

```
0:000> u 00007ff8`1c6fc028
00007ff8`1c6fc028 e9a3401100 jmp     00007ff8`1c8100d0
00007ff8`1c6fc02d 5f      pop     rdi
00007ff8`1c6fc02e 0201    add     al,byte ptr [rcx]
```

```
0:000> u 00007ff8`1c8100d0
```

```
...
```

JIT and NGEN

■ Inlining

- Quite aggressive by default
- UsingMethodImplAttribute andMethodImplOptions
 - NoInlining – prevents the method from getting inlined
 - NoOptimization – turns off optimizations in JIT / NGEN compilation
 - AggressiveInlining – inlines the method whenever possible

■ JIT intrinsics

- Methods provided by the CLR as native code
- E.g. Math.Sin

■ NGEN

- ngen install <assembly>
- Creates _ni file in native image cache
 - %windir%\assembly\NativeImages_v4.0.30319_32
 - Folder per assembly

JIT and NGEN

```
static void Main() {  
    Console.WriteLine(Add(1, 2));  
}  
  
static int Add(int a, int b) {  
    return a + b;  
}
```



0:000> !U 00007FFD85E70090

Normal JIT generated code

Arith.Main()

Begin 00007fffd85e70090, size 14

00007fffd`85e70090 sub rsp,28h

00007fffd`85e70094 mov ecx,3

00007fffd`85e70099 call mscorlib_ni+0xd24780 (00007fffd`e4b44780)
 (System.Console.WriteLine(Int32), mdToken: 000000000600099d)

00007fffd`85e7009e nop

00007fffd`85e7009f add rsp,28h

00007fffd`85e700a3 ret

1 + 2 = 3

Summary

- **Intermediate Language (IL)**
 - Virtual machine with stack-based evaluation
 - Metadata for types and members
 - Static typing all around
 - Rich instruction set
- **Just-in-time compilation (JIT)**
 - Translation of IL to assembler (x86, x64, ARM)
 - Various optimizations such as inlining
 - Use SOS to inspect generated code
 - Can run ahead of time with NGEN